

# Distributed-Reinforcement-Learning-with-Ray-Framework

## Requirements

```
conda env create -f ray.yml
# tensorflow >= 1.10
# ray >= 0.8.5
# ray[tune]
```

## Architecture

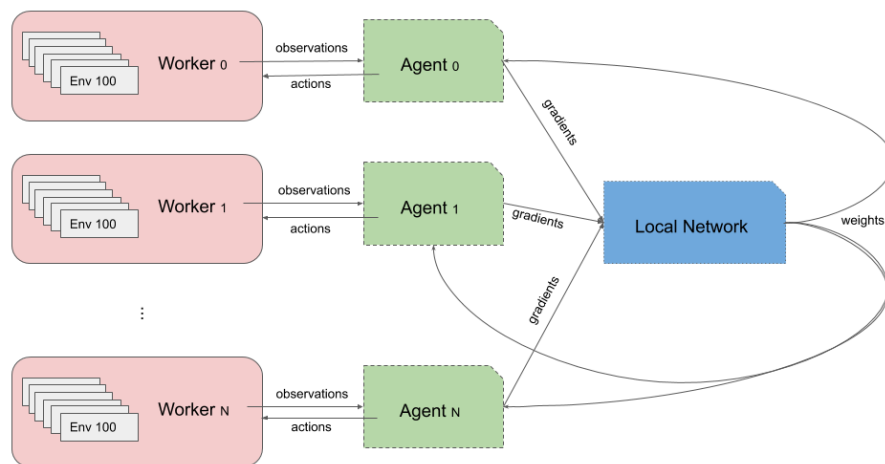


Figure 1: Illustration of distributed architecture.

## Training

```
(ray) $ python main.py --env_name "gym_env" \
--agent_name "default_agent" \
--model_name "ppo" \
--num_envs_per_worker 3 \
--num_workers 3 \
--gpu "0,1,2" \
--num_gpus 3 \
--num_cpus 8 \
--batch_size 10
```

## Custom usage

- Custom environment. Modify `./envs/custom_env.py`

```
@register_env('custom_env')
@ray.remote
class CustomEnvWrapper(object):
    def __init__(self, args):
        self.env = "<custom env>"
        ...

    def reset(self):
        # custom operation
        # 1. ob = self.env.reset()
        # 2. return OBSERVATION(...)

    def step(self, action, **kwargs):
        # 1. process the action
        # 2. env.step
        # #observation = self.env.step(action)
        # 3. return OBSERVATION(...)
```

- Custom models. Modify `models/custom_model.py`

```
@register_model('custom_model')
class CustomModel(object):
    def __init__(self, args):
        ...

    @staticmethod
    def add_args(parser):
        # Add customized models arguments
        ...

    def step(self, *args, **kwargs):
        ...
        # Compute and return gradients

    def choose_actions(self, observation):
        ...
        # Return selected actions (batch)

    def get_weights(self):
        # Return weights
        ...

    def set_weights(self, weights):
```

```

        # Set weights
        ...

    • Custom agents. Modify agents/custom_agent.py
@register_agent("custom_agent")
class DefaultAgent(object):
    def __init__(self, args, model):
        self.model = model
        self.trajectories = []
        ...

    @staticmethod
    def add_args(parser):
        """ Add custom arguments here """
        ...

    def set_weights(self, weights):
        """ Set weights """
        ...

    def send_trajectory(self, trajectory, Rs=None):
        """ Collect trajectories for each worker """
        ...

    def fetch_grads(self, weights):
        """ Compute gradients """
        ...

    • Custom arguments. Modify get_custom_args function in options/custom_options.py.
    • Run the customized model

(ray) $ python main.py --env_name "custom_env" \
      --agent_name "custom_agent" \
      --model_name "custom_model" \
      --num_envs_per_worker 3 \
      --num_workers 3 \
      --gpu "0,1,2" \
      --num_gpus 3 \
      --num_cpus 8 \
      --batch_size 10
      # custom argument options...

```